

Supplement: A four-step Bayesian workflow for improving ecological science

EM Wolkovich, TJ Davies, WD Pearse & M Betancourt

February 21, 2024

A brief review of statistical inference using Bayesian approaches

Robust analyses ely on our inferences being consistent with the underlying truth more often than not. Quantifying this consistency is calibration—analyzing how often a parameter estimate is close to the true value—a critical part of using models for inference. A major problem with traditional (frequentist) approaches in ecology today is their inferences are unpredictable when their foundational assumptions are not met, but ecologists are not usually trained in how to recognize or deal with this.

To better understand our workflow, we provide a very brief overview of some of the fundamentals of Bayesian methods that is inherently incomplete and, by design, not very technical. This section can be skipped for those who feel already well-versed, and can be augmented for those who are new to Bayesian approaches (for example, ???).

Probability is often defined as “the long-term frequency with which something happens.” We would expect, for example that if we tossed a coin 100 times we would see roughly 50 heads. In this case we would say that the probability of tossing a coin and getting a ‘head’ is $\frac{50}{100}$, equivalent to 50% or $\frac{1}{2}$. At the same time we wouldn’t be very surprised if we observed 49 or even 55 heads, although we would be surprised if we saw 99.

This definition of probability—which is the *frequentist* definition—is useful in many situations, but it has a few disadvantages. First, frequentist definitions aren’t very helpful when dealing with unexpected situations. Frequentist probabilities are grounded in repeatable observations, and so understanding these repeatable frequencies is of limited use when trying to make predictions for changing or entirely novel systems. Second, most common frequentist approaches rely on specific assumptions. Using frequentist statistics requires trying to match a model that we have (often just in our heads) of some ecological system to a frequentist method that mostly closely matches the assumptions of our biological model. Given the complexity of ecological data and our uncertainty about the underlying model, frequentist approaches can be especially challenging in ecology.

Bayesian approaches provide a way to build models that can propagate our (un)certainities about what we do—or especially, don’t—know about how ecological systems work. In the Bayesian view, *probability* is used to quantify uncertainty: the higher the probability of a certain interval of values

the less uncertainty we have that the true behavior falls within that interval. Assuming we also have some estimate of our initial uncertainty—usually from knowledge of ecology—to inform a prior distribution (termed below *prior*), then we can apply Bayes’ theorem:

$$probability = \frac{likelihood \cdot prior}{normaliser} \tag{1}$$

to update that prior into a posterior distribution that accounts for the information added by a likelihood (*likelihood* is the same as a frequentist likelihood). Here, **normaliser** is a mathematical constant that makes sure our probability cannot go above 100% or below 0% (statisticians are lazy, and will not ‘give 110%’). This mathematical constant [technically it is the probability of our data; $p(data)$] is a nuisance term that is extremely challenging to estimate (sometimes it is impossible!) and held back the practical use of Bayesian statistics for almost a century because that normaliser could rarely be analytically worked out. But one of the major advantages of Bayesian methods is that the solution to this problem—numerical simulations based on Markov Chain Monte Carlo (MCMC)—provide a huge additional advantage. Now that no analytical solution needs be found, any model that can be written out mathematically can be fit to data, giving the scientist more freedom of model structure.

Which workflow?

Formally, all a ‘workflow’ does is organize various steps together in a systematic fashion, but there are many different workflows depending on what the aim is, which will determine which steps a workflow should include. For example a workflow aimed at calibration could look like an expanded version of our Step 1, where all the steps focus on investigating the assumptions encoded in a given model using simulated data. Or a workflow aimed at inference could expand Step 3, to focus on constructing a posterior, then investigating its model adequacy via several criteria. An inferential workflow can also be extended into a model development workflow. If the model adequacy criteria inform not only that something is inadequate about the current model assumptions but what is inadequate (ideally this happens some in Step 4) then one can use those hints to iterative improve the modeling assumptions. We present in the main text a very simplified model development workflow that combines calibration, inference and some model development, but it is not necessarily appropriate for everyone, depending on their aims.

An example workflow

Our workflow is explained mostly program-agnostically. Though at times we assume a user of **Stan**, a relatively new probabalistic programming language, that interfaces with R, Python, Julia (and more) to write bespoke Bayesian models (?). We focus on **Stan** as its MCMC algorithm (a variant of Hamiltonian Monte Carlo, HMC) is fast and produces specific output to warn of model fit issues (i.e., divergent transitions) in a way other MCMC algorithms do not (e.g. Metropolis-Hastings or Gibbs), but the basic workflow should apply to diverse implementations of Bayesian modeling, and can be extended to other approaches (frequentist, resampling, etc.).

one of the common problems, especially for people who are used to building analyses from black box components, is that transitioning to more flexible techniques is often overwhelming. Even if the ideas are compelling the open-endedness in the implementation can discourage people who experiment, causing them to default back to the status quo.

The benefit of a longer presentation with comprehensive examples is that the reader directly sees how all of these conceptual pieces fall into a practical implementation and how straightforward that implementation can be in the context of a particular application. Generalizing to bespoke applications will still be frustrating, but at least they’re walking in the right direction.

Even before we arrived at the data simulation step, we realized our verbal model did not agree with the statistical model we expected to fit. We planned to fit a simple linear model, but that would assume climate change has been ongoing across all our years and that’s not what most science on anthropogenic warming suggests: it suggests instead a large uptick in warming around 1980 (likely do to effects of aerosols, ?). So instead we developed a ‘hinge’ model to fit the linear regression after 1980 and a mean before 1980. This highlights a reality throughout the workflow: effective model building is about efficient brainstorming. It’s a constant back and forth between asking questions about what we know and what we should know.

Step 1:

Next we simulated data to test our model code. To do this we set the parameters as our first step, then we simulate the data from these set parameters. In simulation, we know the ‘truth’—which here is a our model parameters—and we then compare those with what we estimated to what we started with.

At this step, we actually realized the partial pooling on our intercept was not a good fit to the data (Fig. ??). We assumed from the beginning that we should partially pool on the intercepts because that is how fixing ‘non-independence’ with ‘random effects’ is taught in ecology, but our retrodictive checks made us think twice. We realized that we had data from diverse events across locations and their average values were unlikely to be normally distributed (though they may be if we had focused on a certain more local area, such as temperate sites in Europe. We ended up with partial pooling only on the slopes and not intercepts—a model that made biological sense, but we never would have considered without this workflow.

1 References

Booth, B.B., Dunstone, N.J., Halloran, P.R., Andrews, T. & Bellouin, N. (2012) Aerosols implicated as a prime driver of twentieth-century north atlantic climate variability. *Nature* **484**, 228–32.